

论文索引、编号与官方链接汇总

按所给截图顺序整理；核对日期：2026-07-18

说明：截图中编号 19 对应三篇 CosyVoice 论文，编号 21 对应两篇论文，故分别记为 19.1–19.3 与 21.1–21.2；截图未出现编号 24 和 43，因此未另行虚构对应条目。论文名称采用公开版本中的完整标题，链接优先采用 arXiv 或发布机构官方页面。

编号	论文名	arXiv 号 / 其他号	地址 URL（可点击）
01	Fractal Generative Models	arXiv:2502.17437	arxiv.org/abs/2502.17437
02	Do As I Can, Not As I Say: Grounding Language in Robotic Affordances	arXiv:2204.01691	arxiv.org/abs/2204.01691
03	Flow to the Mode: Mode-Seeking Diffusion Autoencoders for State-of-the-Art Image Tokenization	arXiv:2503.11056	arxiv.org/abs/2503.11056
04	NoProp: Training Neural Networks without Back-propagation or Forward-propagation	arXiv:2503.24322	arxiv.org/abs/2503.24322
05	PC Agent: A Hierarchical Multi-Agent Collaboration Framework for Complex Task Automation on PC	arXiv:2502.14282	arxiv.org/abs/2502.14282
06	5% > 100%: Breaking Performance Shackles of Full Fine-Tuning on Visual Recognition Tasks	arXiv:2408.08345	arxiv.org/abs/2408.08345
07	LoRI: Reducing Cross-Task Interference in Multi-Task Low-Rank Adaptation	arXiv:2504.07448	arxiv.org/abs/2504.07448
08	Building Vision Models upon Heat Conduction	arXiv:2405.16555	arxiv.org/abs/2405.16555
09	Mean Flows for One-step Generative Modeling	arXiv:2505.13447	arxiv.org/abs/2505.13447
10	Exploring Diffusion Transformer Designs via Grafting	arXiv:2506.05340	arxiv.org/abs/2506.05340
11	Localizing Lying in Llama: Understanding Instructed Dishonesty on True-False Questions Through Prompting, Probing, and Patching	arXiv:2311.15131	arxiv.org/abs/2311.15131
12	MUDDFormer: Breaking Residual Bottlenecks in Transformers via Multiway Dynamic Dense Connections	arXiv:2502.12170	arxiv.org/abs/2502.12170

编号	论文名	arXiv 号 / 其他号	地址 URL (可点击)
13	Chain-of-Thought Is Not Explainability	Oxford Martin AIGI Working Paper (2025)	www.oxfordmartin.ox.ac.uk/publications/chain-of-thought-is-not-explainability
14	Multiverse: Your Language Models Secretly Decide How to Parallelize and Merge Generation	arXiv:2506.09991	arxiv.org/abs/2506.09991
15	DiC: Rethinking Conv3x3 Designs in Diffusion Models	arXiv:2501.00603	arxiv.org/abs/2501.00603
16	SepLLM: Accelerate Large Language Models by Compressing One Segment into One Separator	arXiv:2412.12094	arxiv.org/abs/2412.12094
17	Mixture-of-Recursions: Learning Dynamic Recursive Depths for Adaptive Token-Level Computation	arXiv:2507.10524	arxiv.org/abs/2507.10524
18	Native Sparse Attention: Hardware-Aligned and Natively Trainable Sparse Attention	arXiv:2502.11089	arxiv.org/abs/2502.11089
19.1	CosyVoice: A Scalable Multilingual Zero-shot Text-to-speech Synthesizer based on Supervised Semantic Tokens	arXiv:2407.05407	arxiv.org/abs/2407.05407
19.2	CosyVoice 2: Scalable Streaming Speech Synthesis with Large Language Models	arXiv:2412.10117	arxiv.org/abs/2412.10117
19.3	CosyVoice 3: Towards In-the-wild Speech Generation via Scaling-up and Post-training	arXiv:2505.17589	arxiv.org/abs/2505.17589
20	OverLoCK: An Overview-first-Look-Closely-next ConvNet with Context-Mixing Dynamic Kernels	arXiv:2502.20087	arxiv.org/abs/2502.20087
21.1	Cambrian-1: A Fully Open, Vision-Centric Exploration of Multimodal LLMs	arXiv:2406.16860	arxiv.org/abs/2406.16860
21.2	Vision Language Models are Biased	arXiv:2505.23941	arxiv.org/abs/2505.23941
22	BriLLM: Brain-Inspired Large Language Model	arXiv:2503.11299	arxiv.org/abs/2503.11299
23	Scalable Diffusion Models with Transformers	arXiv:2212.09748	arxiv.org/abs/2212.09748
25	DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning	arXiv:2501.12948	arxiv.org/abs/2501.12948

编号	论文名	arXiv 号 / 其他号	地址 URL (可点击)
26	Why Language Models Hallucinate	arXiv:2509.04664	arxiv.org/abs/2509.04664
27	Optical Generative Models	arXiv:2410.17970	arxiv.org/abs/2410.17970
28	DeepSeek-OCR: Contexts Optical Compression	arXiv:2510.18234	arxiv.org/abs/2510.18234
29	FlyLoRA: Boosting Task Decoupling and Parameter Efficiency via Implicit Rank-Wise Mixture-of-Experts	arXiv:2510.08396	arxiv.org/abs/2510.08396
30	Back to Basics: Let Denoising Generative Models Denoise	arXiv:2511.13720	arxiv.org/abs/2511.13720
31	ARC Is a Vision Problem!	arXiv:2511.14761	arxiv.org/abs/2511.14761
32	Kimi K2.5: Visual Agentic Intelligence	arXiv:2602.02276	arxiv.org/abs/2602.02276
33	Attention Residuals	arXiv:2603.15031	arxiv.org/abs/2603.15031
34	Toy Models of Superposition	arXiv:2209.10652	arxiv.org/abs/2209.10652
35	Obscure but Effective: Classical Chinese Jailbreak Prompt Optimization via Bio-Inspired Search	arXiv:2602.22983	arxiv.org/abs/2602.22983
36	Vision Transformers Need More Than Registers	arXiv:2602.22394	arxiv.org/abs/2602.22394
37	Artificial Intelligence Index Report 2026	arXiv:2606.15708 / Stanford HAI Report	hai.stanford.edu/ai-index/2026-ai-index-report
38	DeepSeek LLM: Scaling Open-Source Language Models with Longtermism	arXiv:2401.02954	arxiv.org/abs/2401.02954
39	Continuous Thought Machines	arXiv:2505.05522	arxiv.org/abs/2505.05522
40	Improved Baselines with Representation Autoencoders	arXiv:2605.18324	arxiv.org/abs/2605.18324
41	Image Generators are Generalist Vision Learners	arXiv:2604.20329	arxiv.org/abs/2604.20329
42	ChordEdit: One-Step Low-Energy Transport for Image Editing	arXiv:2602.19083	arxiv.org/abs/2602.19083
44	Mapping Networks	arXiv:2602.19134	arxiv.org/abs/2602.19134
45	DSpark: Confidence-Scheduled Speculative Decoding with Semi-Autoregressive Generation	arXiv:2607.05147	arxiv.org/abs/2607.05147